

install kubernetes 1.7.5 on ubuntu 1604 (virtualbox vm)

写在前面：

K8s环境搭建是所有练习的第一步，也是所有练习的基础。

重点：

理解K8s的应用场景，基本组件和原理是本次练习的重点，也是为后面几个练习做准备，所以完成练习只是一个途径，更多的是要对K8s有一个比较完整的理解。下面会列出几个供学习和参考的网址，大家可以提前准备一下：

1. K8s官网地址：<https://kubernetes.io/docs/home/>
2. Jimmysong的K8s handbook：<https://jimmysong.io/kubernetes-handbook/>
3. Docker官网地址：<https://www.docker.com/>

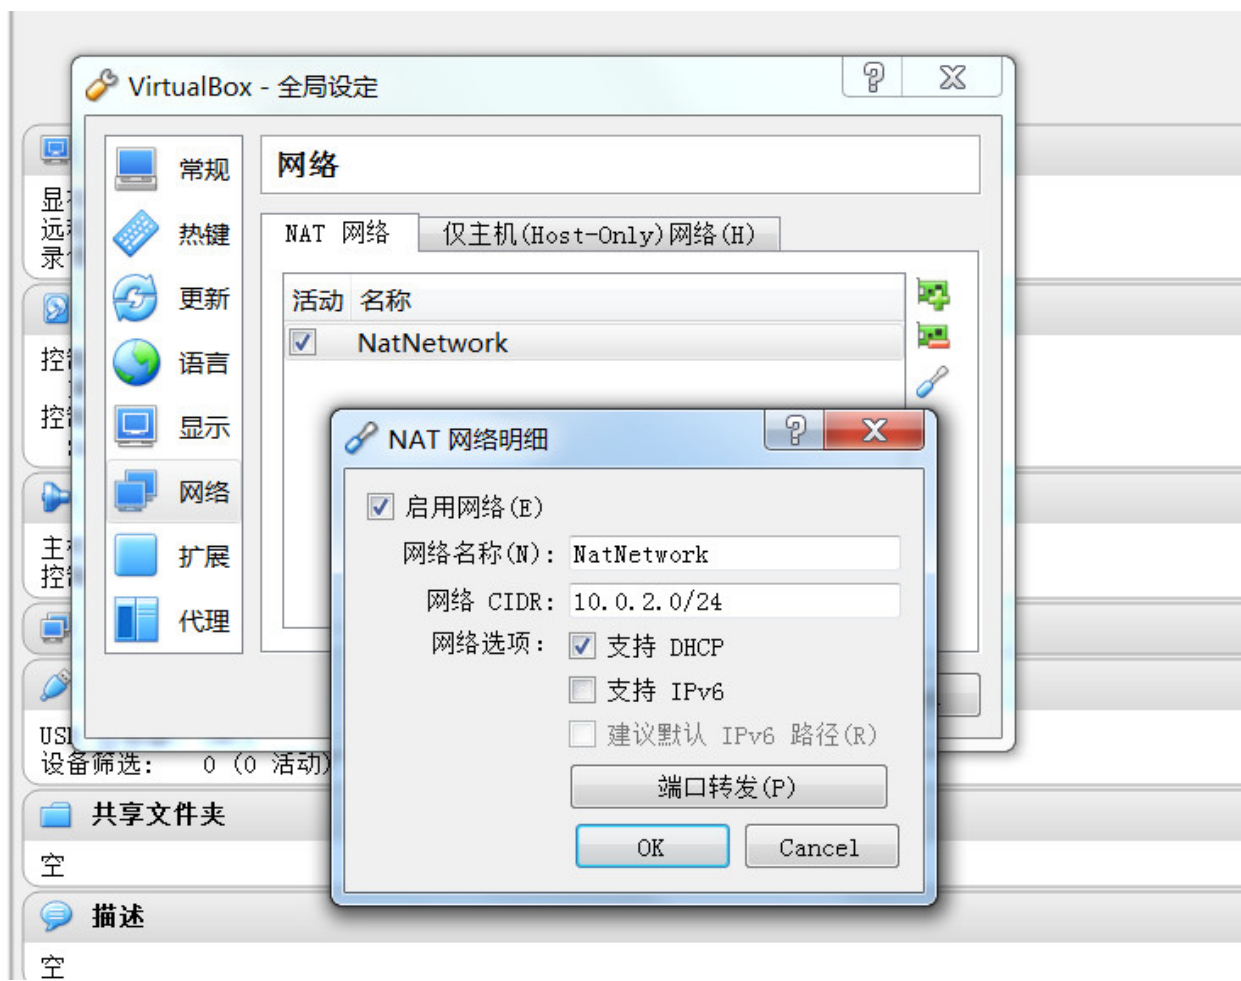
难点：

第一个练习本身并不难，更多的问题可能出现在虚拟机环境的准备上。下面列出几个可能碰到的问题。

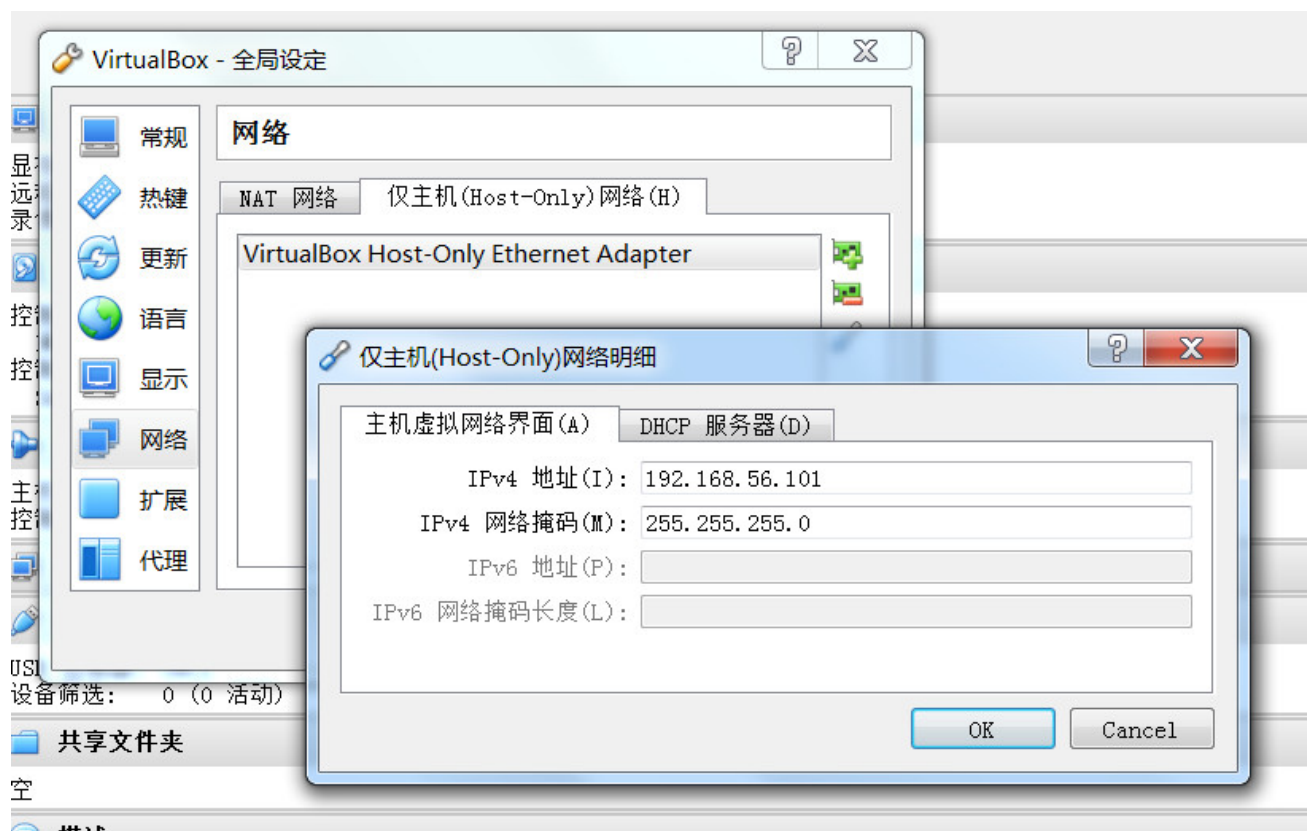
坑1：网络配置

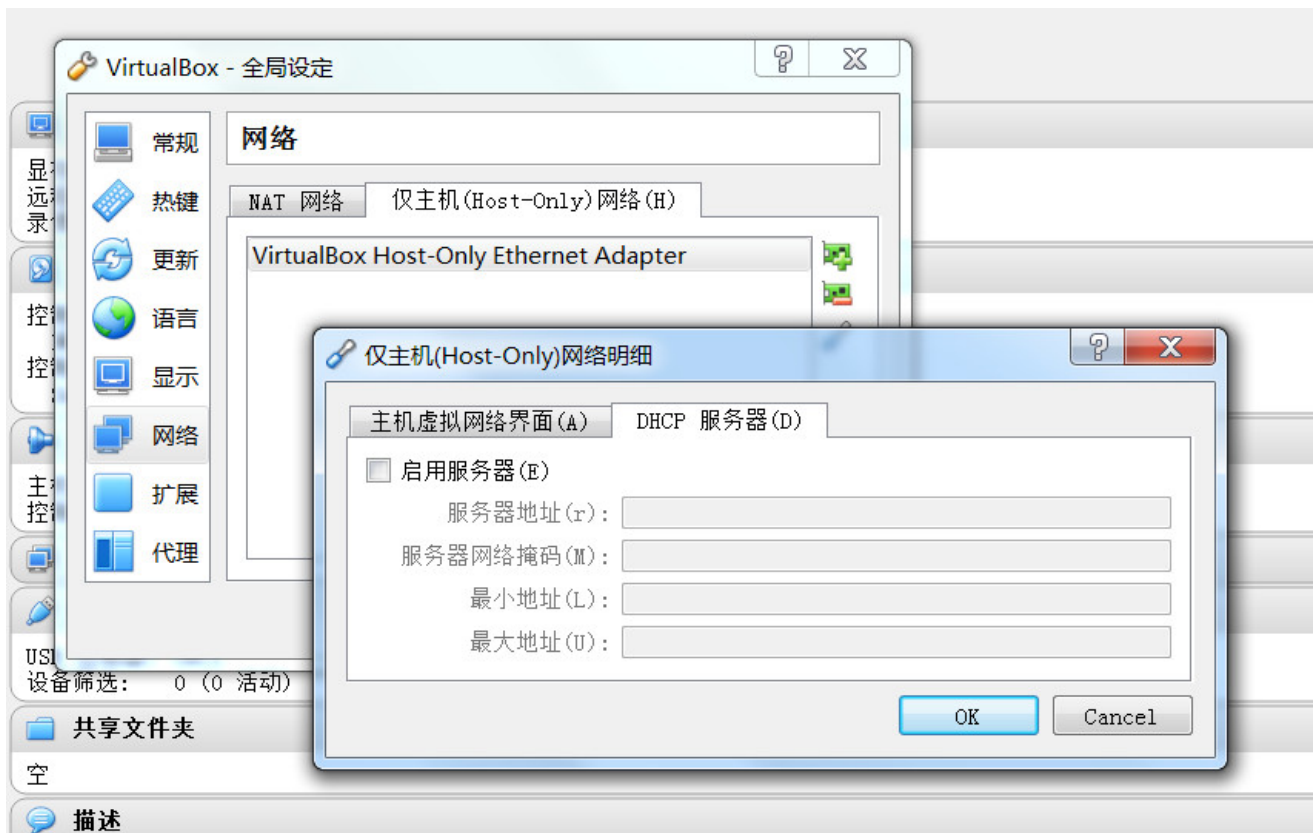
实验一本身不需要链接互联网，但是后续的实验需要进行联网下载镜像，所以建议在第一个实验时就配置好相关网络。配置过程如下：

在VirtualBox界面上点击"管理"-">"全局设定"，在弹出框中选择"网络"-">"NAT网络"，新建一个"NAT网络"，具体如下：

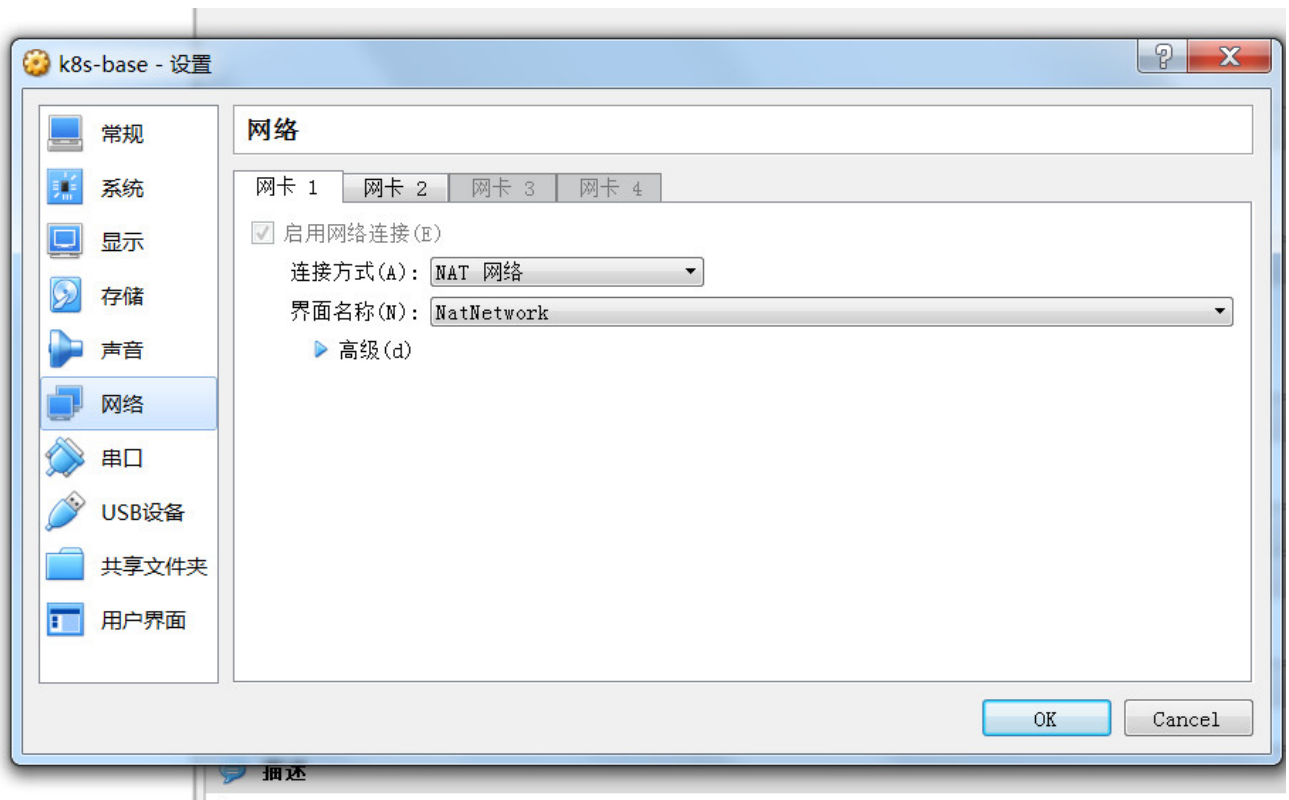


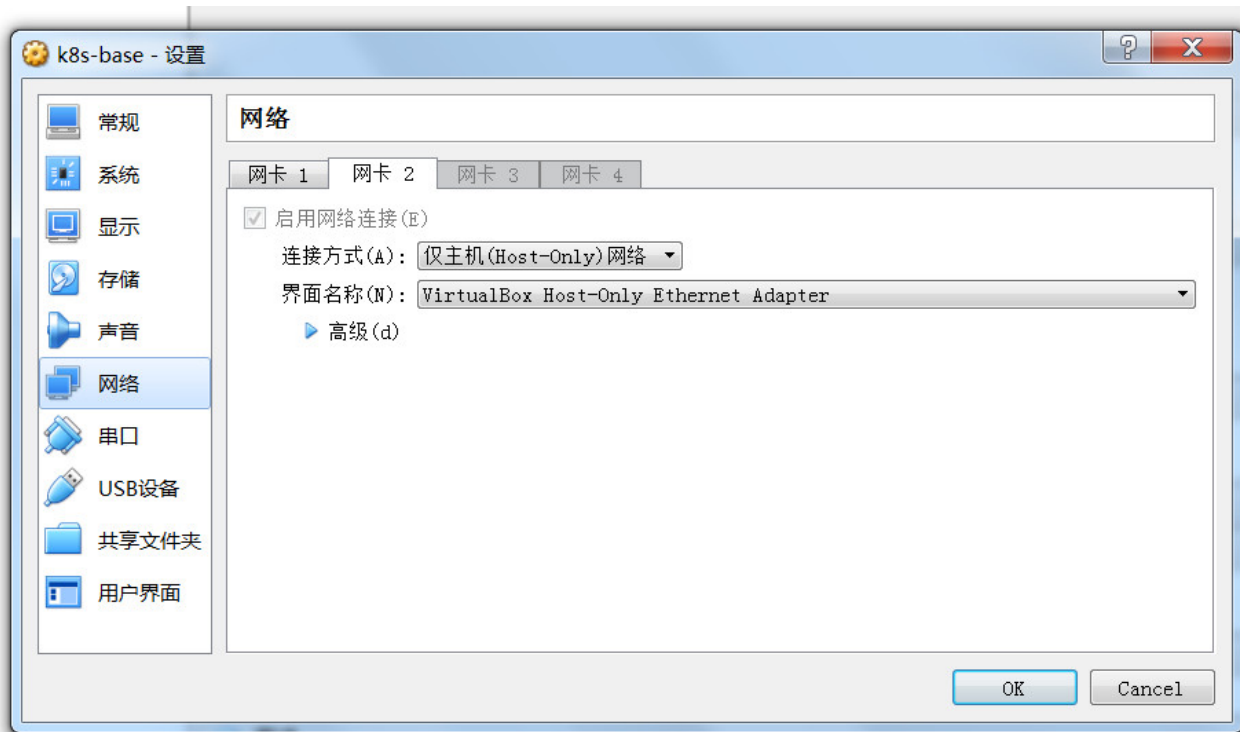
选择"仅主机(Host-Only)网络(H)", 设置虚拟机内部网卡地址, 如下:



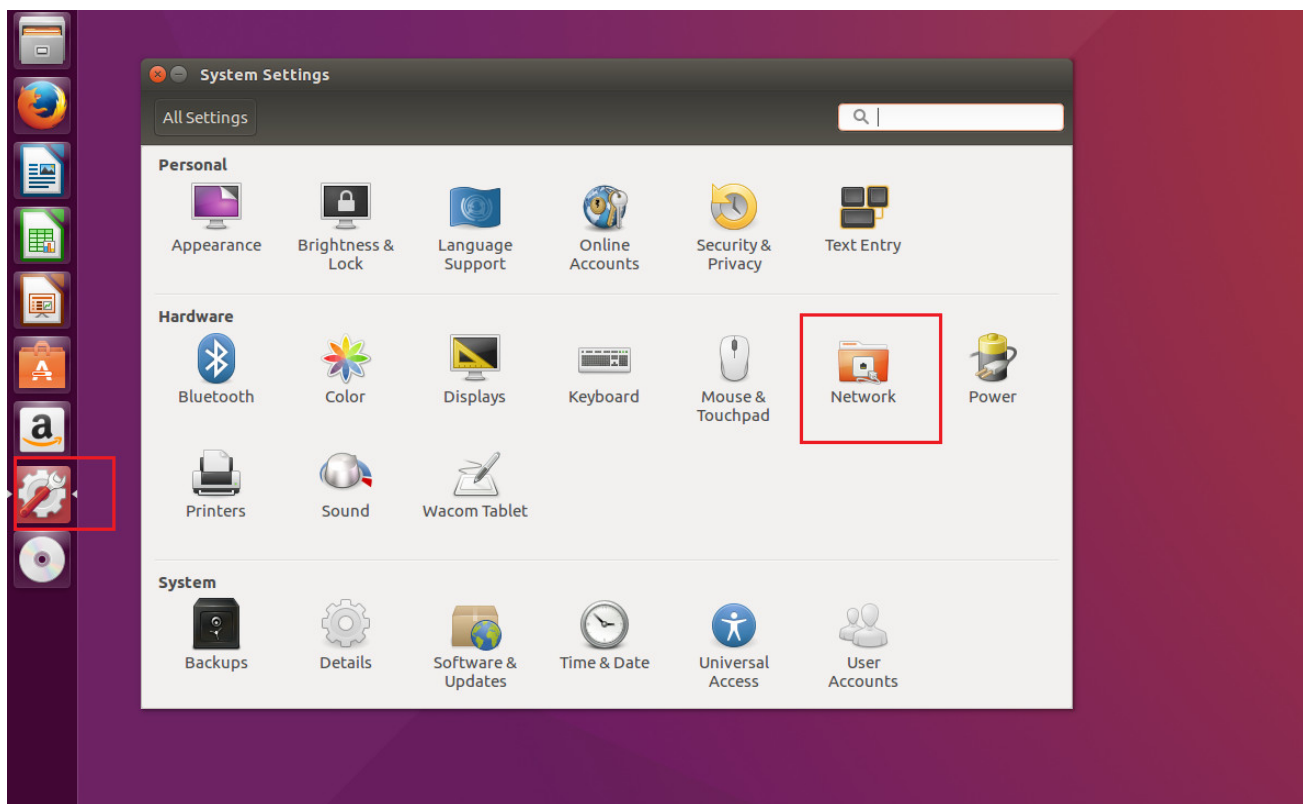


右键虚拟机，选择"设置"-"网络"，新建"NAT 网络"和"仅主机(Host-Only)网络(H)"，如下：





运行虚拟机，对网卡进行设置，如下：



Network

All SettingsNetwork

Airplane ModeOFF

Wired

Wired

Network proxy

Wired

Connected - 1000 Mb/s

Hardware Address 08:00:27:5A:93:BB

IPv4 Address 10.0.2.5

IPv6 Address fe80::90f3:92cf:a28f:a62b

Default Route 10.0.2.1

DNS 10.122.1.82 10.122.2.82

Options...

Editing Wired connection 2

Connection name: Wired connection 2

General

Ethernet

802.1x Security

DCB

IPv4 Settings

IPv6 Settings

Method: Automatic (DHCP)

Addresses

Address	Netmask	Gateway	
			Add
			Delete

Additional DNS servers:

Additional search domains:

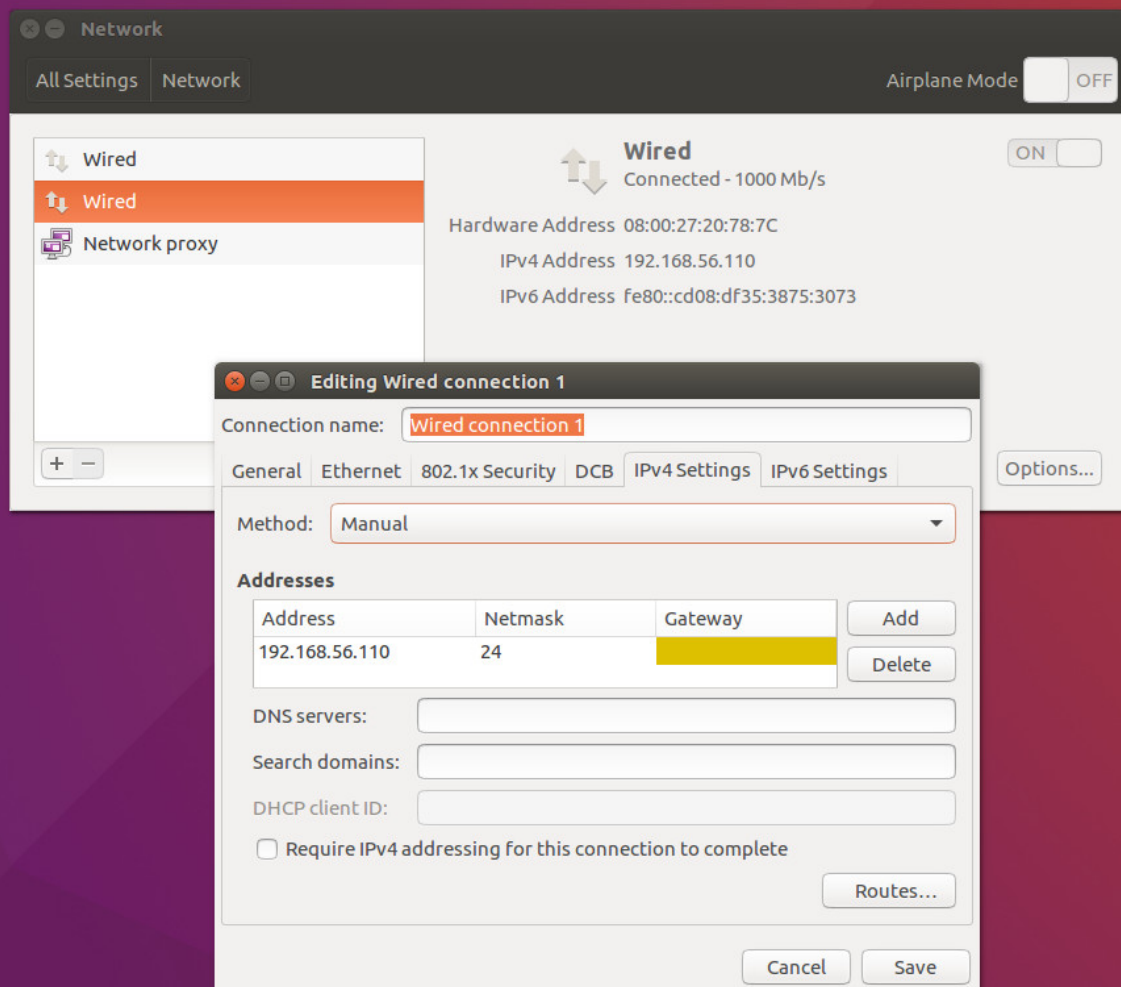
DHCP client ID:

☐ Require IPv4 addressing for this connection to complete

Routes...

Cancel

Save



重启虚拟机，使用route命令，检查网卡设置是否正确：

```
svw@master: ~
svw@master:~$ route
Kernel IP routing table
Destination    Gateway         Genmask         Flags Metric Ref    Use Iface
default        10.0.2.1        0.0.0.0         UG    100    0      0 enp0s3
10.0.2.0       *               255.255.255.0   U     100    0      0 enp0s3
link-local     *               255.255.0.0     U     1000   0      0 enp0s8
172.17.0.0     *               255.255.0.0     U     0      0      0 docker0
192.168.56.0   *               255.255.255.0   U     100    0      0 enp0s8
svw@master:~$
```

使用ping命令，检查是否能够链接外网。

坑2：hosts文件配置


配置hosts文件时，要保留localhost的配置，否则会造成master或者node节点启动失败，如下：


```
svw@master:~$ cat /etc/hosts
127.0.0.1      master localhost
127.0.1.1      svw-poc

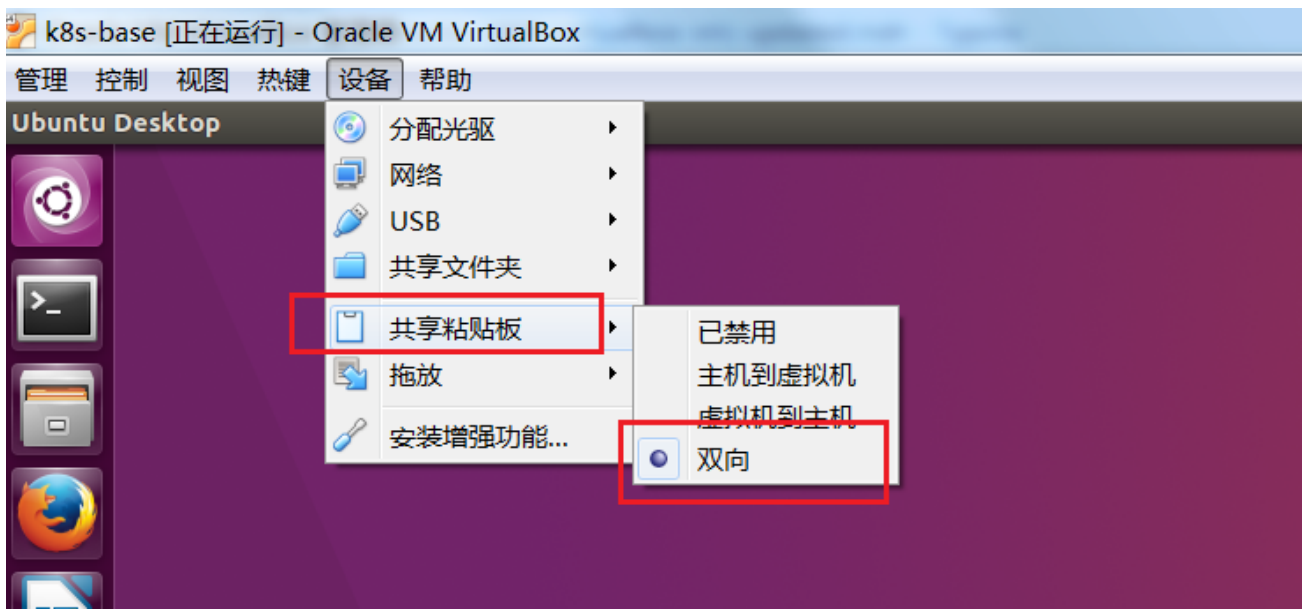
# The following lines are desirable for IPv6 capable hosts
::1           ip6-localhost ip6-loopback
fe00::0       ip6-localnet
ff00::0       ip6-mcastprefix
ff02::1       ip6-allnodes
ff02::2       ip6-allrouters
```

建议：

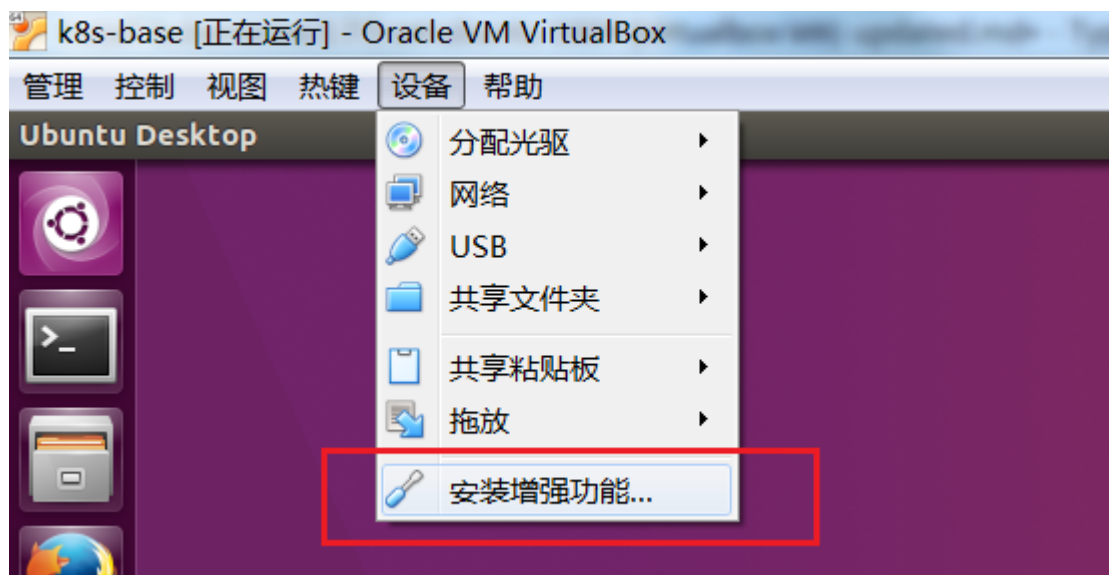
建议1：创建一个基础虚拟机，安装基本功能和设置，其它的虚拟机从这个基础虚拟机复制后，再进行练习，如下：




K8s-base中安装与宿主机进行同步拷贝，粘贴功能：



安装增强功能：



点击“Run”，输入管理员密码“svw”，并等待安装完成：



宿主机环境要求

- VirtualBox 5.1.22以上版本
- 虚拟机网络启用virtualbox host-only模式，网卡IP地址设定为192.168.56.1

虚拟机模板说明

操作系统版本：ubuntu 16.04

账号：svw/svw

虚拟机网卡：单网卡，host-only模式；注意务必不要启用NAT网卡，那会造成flannel网络组件获取地址错误。

虚拟机内已安装组件清单：ssh, curl, docker, kublet, kubeadm, kubectl

安装master节点

从模板复制虚拟机

基于提供的模板虚拟机复制出一个新的虚拟机（注意务必选中 重新初始化所有的网卡MAC地址）

设置静态IP地址

```
sudo vi /etc/network/interfaces
```

将以下内容添加到文件中保存退出。

```
# The primary network interface
auto enp0s3

iface enp0s3 inet static
address 192.168.56.110
netmask 255.255.255.0
gateway 192.168.56.1
```

更改主机名

```
sudo vi /etc/hostname
```

将主机名改为master，保存退出

更改hosts文件

```
sudo vi /etc/hosts
```

将127.0.0.1对应的主机名改为master，保存退出

重启服务器

```
sudo init 6
```

导入镜像

将K8s所需要的镜像(google_containers.tar)上转至服务器目录/home/swap (通过winscp上传)

```
sudo docker load -i /home/swap/google_containers.tar
```

查看下导入的镜像列表

```
svw@master:~$ sudo docker images
```

REPOSITORY	TAG	IMAGE ID
gcr.io/google_containers/kube-apiserver-amd64	v1.7.5	33cecf00bef1
3 weeks ago	186.1 MB	
gcr.io/google_containers/kube-controller-manager-amd64	v1.7.5	4d581c01252d
3 weeks ago	138 MB	
gcr.io/google_containers/kube-scheduler-amd64	v1.7.5	2aad6febbb16
3 weeks ago	77.2 MB	
gcr.io/google_containers/kube-proxy-amd64	v1.7.5	1314eb3ac430
3 weeks ago	114.7 MB	
quay.io/coreos/flannel	v0.8.0-amd64	9db3bab8c19e
11 weeks ago	50.73 MB	
gcr.io/google_containers/k8s-dns-sidecar-amd64	1.14.4	38bac66034a6
3 months ago	41.82 MB	
gcr.io/google_containers/k8s-dns-kube-dns-amd64	1.14.4	a8e00546bcf3
3 months ago	49.39 MB	
gcr.io/google_containers/k8s-dns-dnsmasq-nanny-amd64	1.14.4	f7f45b9cb733
3 months ago	41.42 MB	
gcr.io/google_containers/etcd-amd64	3.0.17	243830dae7dd
7 months ago	168.9 MB	
gcr.io/google_containers/pause-amd64	3.0	99e59f495ffa
17 months ago	746.9 kB	

安装master上的K8s组件

```
sudo kubeadm init --kubernetes-version v1.7.5 --apiserver-advertise-address 192.168.56.110 --pod-network-cidr 10.244.0.0/16
```

其中，`--apiserver-advertise-address` 对应的值为master节点的本机IP地址，`--pod-network-cidr` 对应的值为k8s集群内部的网络段。

运行成功后会有以下内容打印出

```
svw@master:~$ kubeadm init --kubernetes-version v1.7.5 --apiserver-advertise-address 192.168.56.110 --pod-network-cidr 10.244.0.0/16
[kubeadm] WARNING: kubeadm is in beta, please do not use it for production clusters.
[init] Using Kubernetes version: v1.7.5
[init] Using Authorization modes: [Node RBAC]
[preflight] Running pre-flight checks
[preflight] Some fatal errors occurred:
        user is not running as root
[preflight] If you know what you are doing, you can skip pre-flight checks with `--skip-preflight-checks`
svw@master:~$ sudo kubeadm init --kubernetes-version v1.7.5 --apiserver-advertise-address 192.168.56.110 --pod-network-cidr 10.244.0.0/16
[kubeadm] WARNING: kubeadm is in beta, please do not use it for production clusters.
[init] Using Kubernetes version: v1.7.5
[init] Using Authorization modes: [Node RBAC]
[preflight] Running pre-flight checks
[kubeadm] WARNING: starting in 1.8, tokens expire after 24 hours by default (if you require a
```

```

non-expiring token use --token-ttl 0)
[certificates] Generated CA certificate and key.
[certificates] Generated API server certificate and key.
[certificates] API Server serving cert is signed for DNS names [master kubernet
kubernetes.default kubernetes.default.svc kubernetes.default.svc.cluster.local] and IPs
[10.96.0.1 192.168.56.110]
[certificates] Generated API server kubelet client certificate and key.
[certificates] Generated service account token signing key and public key.
[certificates] Generated front-proxy CA certificate and key.
[certificates] Generated front-proxy client certificate and key.
[certificates] Valid certificates and keys now exist in "/etc/kubernetes/pki"
[kubeconfig] Wrote KubeConfig file to disk: "/etc/kubernetes/kubelet.conf"
[kubeconfig] Wrote KubeConfig file to disk: "/etc/kubernetes/controller-manager.conf"
[kubeconfig] Wrote KubeConfig file to disk: "/etc/kubernetes/scheduler.conf"
[kubeconfig] Wrote KubeConfig file to disk: "/etc/kubernetes/admin.conf"
[apiclient] Created API client, waiting for the control plane to become ready
[apiclient] All control plane components are healthy after 36.500720 seconds
[token] Using token: e99d39.674b685c4270a1d7
[apiconfig] Created RBAC rules
[addons] Applied essential addon: kube-proxy
[addons] Applied essential addon: kube-dns

```

Your Kubernetes master has initialized successfully!

To start using your cluster, you need to run (as a regular user):

```

mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config

```

You should now deploy a pod network to the cluster.

Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:
<http://kubernetes.io/docs/admin/addons/>

You can now join any number of machines by running the following on each node
as root:

```

kubeadm join --token e99d39.674b685c4270a1d7 192.168.56.110:6443

```

注意：最后一句需要保存好，未来安装新的node节点时都会用到。

配置kubectl，在svw用户环境下逐条执行以下命令

```

mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config

```

执行kubectl检查各组件的状态，正常情况下会有以下内容输出

```

svw@master:~$ kubectl get all --all-namespaces

```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
kube-system	po/etcd-master	1/1	Running	0	2m

kube-system	po/kube-apiserver-master	1/1	Running	0	2m
kube-system	po/kube-controller-manager-master	1/1	Running	0	2m
kube-system	po/kube-dns-2425271678-hmblq	0/3	Pending	0	3m
kube-system	po/kube-proxy-cp11m	1/1	Running	0	3m
kube-system	po/kube-scheduler-master	1/1	Running	0	2m

NAMESPACE	NAME	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
default	svc/kubernetes	10.96.0.1	<none>	443/TCP	3m
kube-system	svc/kube-dns	10.96.0.10	<none>	53/UDP,53/TCP	3m

NAMESPACE	NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
kube-system	deploy/kube-dns	1	1	1	0	3m

NAMESPACE	NAME	DESIRED	CURRENT	READY	AGE
kube-system	rs/kube-dns-2425271678	1	1	0	3m

注：dns组件状态为pending是因为网络组件还未安装完成，属正常情况。

安装网络组件flannel

将kube-flannel.yml、kube-flannel-rbac.yml上传至/home/swv

执行安装命令

```
kubectl apply -f kube-flannel-rbac.yml
kubectl apply -f kube-flannel.yml
```

执行完成后等待两分钟后会有如下输出，master节点安装完成。

```
swv@master:~$ kubectl get all --all-namespaces
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
kube-system	po/etcd-master	1/1	Running	0	2m
kube-system	po/kube-apiserver-master	1/1	Running	0	2m
kube-system	po/kube-controller-manager-master	1/1	Running	0	2m
kube-system	po/kube-dns-2425271678-hmblq	3/3	Running	0	7m
kube-system	po/kube-flannel-ds-rqp7j	2/2	Running	0	42s
kube-system	po/kube-proxy-cp11m	1/1	Running	0	7m
kube-system	po/kube-scheduler-master	1/1	Running	0	2m

NAMESPACE	NAME	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
default	svc/kubernetes	10.96.0.1	<none>	443/TCP	7m
kube-system	svc/kube-dns	10.96.0.10	<none>	53/UDP,53/TCP	7m

NAMESPACE	NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
kube-system	deploy/kube-dns	1	1	1	1	7m

NAMESPACE	NAME	DESIRED	CURRENT	READY	AGE
kube-system	rs/kube-dns-2425271678	1	1	1	7m

安装node1节点

从模板复制虚拟机

基于提供的模板虚拟机复制出一个新的虚拟机（注意务必选中 重新初始化所有的网卡MAC地址）

设置静态IP地址

```
sudo vi /etc/network/interfaces
```

将以下内容添加到文件中保存退出。

```
# The primary network interface
auto enp0s3

iface enp0s3 inet static
address 192.168.56.111
netmask 255.255.255.0
gateway 192.168.56.1
```

更改主机名

```
sudo vi /etc/hostname
```

将主机名改为node1，保存退出

更改hosts文件

```
sudo vi /etc/hosts
```

将127.0.0.1对应的主机名改为node1，保存退出

重启服务器

```
sudo init 6
```

导入镜像

将K8s所需要的镜像(google_containers.tar)上转至服务器目录/home/svw（通过winscp上传）

```
sudo docker load -i /home/svw/google_containers.tar
```

将node1加入到master节点

执行以下命令将node1加入到master节点

```
sudo kubeadm join --token e99d39.674b685c4270a1d7 192.168.56.110:6443
```

执行成功会得到以下输出

```
svw@node1:~$ sudo kubeadm join --token e99d39.674b685c4270a1d7 192.168.56.110:6443
```

```
[kubeadm] WARNING: kubeadm is in beta, please do not use it for production clusters.
[preflight] Running pre-flight checks
[discovery] Trying to connect to API Server "192.168.56.110:6443"
[discovery] Created cluster-info discovery client, requesting info from
"https://192.168.56.110:6443"
[discovery] Cluster info signature and contents are valid, will use API Server
"https://192.168.56.110:6443"
[discovery] Successfully established connection with API Server "192.168.56.110:6443"
[bootstrap] Detected server version: v1.7.5
[bootstrap] The server supports the Certificates API (certificates.k8s.io/v1beta1)
[csr] Created API client to obtain unique certificate for this node, generating keys and
certificate signing request
[csr] Received signed certificate from the API server, generating KubeConfig...
[kubeconfig] Wrote KubeConfig file to disk: "/etc/kubernetes/kubelet.conf"
```

Node join complete:

- * Certificate signing request sent to master and response received.
- * Kubelet informed of new secure connection details.

Run 'kubectl get nodes' on the master to see this machine join.

master节点上查看整个群集的状态

```
svw@master:~$ kubectl get nodes
NAME        STATUS    AGE         VERSION
master      Ready     1h          v1.7.5
node1       Ready     31s         v1.7.5
svw@master:~$ kubectl get all --all-namespaces
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
kube-system	po/etcd-master	1/1	Running	0	1h
kube-system	po/kube-apiserver-master	1/1	Running	0	1h
kube-system	po/kube-controller-manager-master	1/1	Running	0	1h
kube-system	po/kube-dns-2425271678-hmblq	3/3	Running	0	1h
kube-system	po/kube-flannel-ds-55kfm	2/2	Running	0	34s
kube-system	po/kube-flannel-ds-rqp7j	2/2	Running	0	1h
kube-system	po/kube-proxy-04fbb	1/1	Running	0	34s
kube-system	po/kube-proxy-cp11m	1/1	Running	0	1h
kube-system	po/kube-scheduler-master	1/1	Running	0	1h

NAMESPACE	NAME	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
default	svc/kubernetes	10.96.0.1	<none>	443/TCP	1h
kube-system	svc/kube-dns	10.96.0.10	<none>	53/UDP,53/TCP	1h

NAMESPACE	NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
kube-system	deploy/kube-dns	1	1	1	1	1h

NAMESPACE	NAME	DESIRED	CURRENT	READY	AGE
kube-system	rs/kube-dns-2425271678	1	1	1	1h